

Product Requirements Document (PRD)

Hey Pebble AI Concierge (MVP)

Version: 1.0 **Date:** July 15, 2025 **Author:** Patrick

1. Introduction & Purpose

This document outlines the requirements for the Minimum Viable Product (MVP) of the "Hey Pebble AI Concierge," an AI-powered assistant designed to enhance the event experience on the Pebble platform. The primary goal of the MVP is to validate the core value proposition: **automating attendee support to reduce the load on human staff and improve attendee satisfaction.**

2. Problem Statement

Event organizers and their volunteer staff are often overwhelmed by a high volume of repetitive questions from attendees regarding schedules, locations, and policies. This leads to long wait times at info booths, staff burnout, and a diminished attendee experience. Attendees struggle to get quick, consistent answers to their questions, causing frustration and wasted time.

3. User Personas & Stories

- **As an Attendee, I want to...**
 - ...ask a question in plain language (text or voice) and get an instant answer.
 - ...find out when and where a specific artist is playing.
 - ...find vendors that sell a specific item (e.g., "coffee") or meet a dietary need (e.g., "vegan").
 - ...get clarity on event policies (e.g., "Can I bring a water bottle?").
 - ...be directed to human help if my question is complex or an emergency.
- **As an Event Organizer, I want to...**
 - ...provide a self-service tool to answer the top 80% of attendee questions automatically.
 - ...ensure urgent issues are flagged and escalated to my on-site staff immediately.
 - ...view analytics on what questions are being asked most frequently to improve future event planning.

4. MVP Features & Scope

In-Scope for MVP:

- **Conversational Interface:** A web-based chat widget that can be embedded on an event's mobile website or accessed via a QR code.
- **Intent Handling:** The assistant must reliably handle the following intents:
 1. **Schedule Lookups:** Retrieve artist/session times and locations.
 2. **Vendor Lookups:** Retrieve vendor locations and primary offerings.
 3. **FAQ Retrieval:** Answer questions from a pre-loaded knowledge base.
 4. **Emergency Escalation:** Detect urgent keywords and direct users to on-site staff.
- **Data Backend:** Integration with a Supabase or Airtable database for storing FAQs, vendor lists, and schedules.
- **Basic Answer Cards:** Responses will be primarily text-based. For location-based answers, the text will describe the location (e.g., "Booth B12, near the East Entrance").
- **Basic Analytics:** The system will log every query and its classified intent. An admin dashboard will show a simple table of the most frequent queries.
- **Multilingual Support (2 Languages):** The assistant will support English and Spanish.

Out-of-Scope for MVP:

- Interactive, clickable maps.
- Voice-out (Text-to-Speech) responses.
- Self-service admin panel for organizers to update data.
- Proactive notifications or suggestions.
- Integration with ticketing or payment systems.

5. Technical Requirements

- **Frontend:** A responsive chat widget built in React (or using Voiceflow for rapid prototyping).
- **Backend:** A Node.js server (or serverless functions) to handle API requests.
- **LLM Provider:** OpenAI (GPT-4) for intent classification and natural language understanding.
- **Database:** Supabase or Airtable.
- **Hosting:** Vercel, Render, or Supabase Edge Functions.

6. Success Metrics

The success of the MVP will be measured against the following targets at a pilot event:

- **Support Deflection Rate:** Achieve a >60% reduction in questions asked at the main information booth compared to a similar past event.
- **Median Time-to-Answer:** The AI's median response time must be under 3 seconds.
- **User Helpfulness Rating:** A simple post-interaction "Was this helpful? (👍/👎)" survey should achieve a >90% positive rating.
- **FAQ Coverage:** The AI should successfully answer >80% of incoming queries without needing to escalate (excluding emergencies).